

SAMPLE CHAPTER · FREE PREVIEW

Chapter 11: Prompt Engineering & Advanced Techniques

From Building Conversational AI with LLMs and Agents

Alexander Apartsin, Ph.D. & Yehudit Aperstein, Ph.D.

Fifth Edition · 2026

3 sections assembled with full code highlighting and rendered math.

Like what you read? The full book has 39 chapters and 22 appendices, covering everything from PyTorch foundations to production AI agents. Available at **llmbook.apartsin.com** and on Amazon.

Chapter 11: Prompt Engineering & Advanced Techniques

"The art of prompting is less about telling a machine what to do and more about learning what it already knows how to do, if only you ask the right way."

—  **Prompt**, Silver-Tongued AI Agent

 Prompt Engineering and Advanced Techniques chapter illustration

Figure 11.0.1: *The same model can be brilliant or baffling depending on how you ask. Prompt engineering is the art of asking the right way.*

Chapter Overview

Prompting is programming with natural language. Every interaction with a large language model begins with a prompt, and the quality of that prompt determines the quality of the output. Yet most practitioners treat prompt engineering as an ad hoc trial-and-error process rather than a systematic discipline. This chapter changes that by presenting prompt engineering as a structured craft with well-defined techniques, [measurable outcomes](#), and principled optimization strategies.

We begin with the **foundational techniques**: zero-shot and few-shot prompting, role assignment, system prompt design, and template construction. Next, we explore **reasoning strategies** that unlock the model's ability to solve complex problems: chain-of-thought prompting, self-consistency, tree-of-thought exploration, and the [ReAct framework that interleaves reasoning with action](#). The third section covers **advanced patterns** including self-reflection loops, meta-prompting, prompt chaining, and automated prompt optimization with DSPy. Finally, we address the critical topics of **prompt security and optimization**: [injection attacks](#), defense

strategies, structured output enforcement, prompt compression, and systematic testing.

By the end of this chapter, you will have a practical toolkit for designing, composing, and securing prompts across a wide range of applications, from simple classification tasks to complex [multi-step reasoning pipelines](#).

Big Picture

Prompt engineering is the most accessible and often the most cost-effective way to improve LLM output quality. The techniques here, including few-shot prompting, chain-of-thought, and structured output generation, apply directly to RAG systems (Chapter 20), agents (Chapter 22), and evaluation (Chapter 29).

LEARNING OBJECTIVES

- Design effective zero-shot, few-shot, and role-based prompts with measurable quality improvements
- Construct system prompts and prompt templates with variable injection for production applications
- Implement chain-of-thought, self-consistency, and tree-of-thought reasoning strategies
- Apply the ReAct framework to interleave reasoning with [external tool use](#)
- Build self-reflection and iterative refinement loops that improve output quality across multiple passes
- Use meta-prompting and prompt chaining to decompose complex tasks into manageable sub-tasks
- Identify and defend against prompt injection attacks (direct, indirect, and jailbreak variants)
- Enforce structured output with JSON mode, Pydantic models, and the [Instructor library](#) (covered in depth in [Section 10.2](#); this chapter addresses the security and reliability dimensions of structured output)
- Apply automated prompt optimization using DSPy, OPRO, and prompt compression techniques like LLMingua

- Implement context engineering with MCP (Model Context Protocol) for dynamic context assembly in production applications
- Design [prompt testing suites](#) with regression tests, A/B experiments, and version control

Prerequisites

- [Chapter 05: Decoding Strategies](#) (temperature, sampling, how generation works)
- [Chapter 10: Working with LLM APIs](#) (API calls, message formats, parameter tuning)
- Basic [Python programming](#) and familiarity with the OpenAI or Anthropic client libraries
- Conceptual understanding of how [transformer models](#) process and generate text

Foundational Prompt Design

Zero-shot, few-shot, role prompting, and the art of clear instruction

The quality of a question determines the quality of an answer. This has always been true; now we have machines to prove it at scale.

—  **Prompt**, Socratic AI Agent

Big Picture

Why does prompting matter? A large language model is a powerful text completion engine, but its output is entirely shaped by its input. The same model that produces vague, rambling answers to a poorly worded question can deliver precise, structured, expert-level responses when given a well-designed prompt. Prompt engineering is the discipline of systematically crafting inputs to maximize output quality. This section covers the foundational techniques that every practitioner needs: zero-shot prompting, few-shot learning, role assignment, system prompt architecture, and template design. These techniques form the building blocks for every advanced strategy covered later in this chapter. The [in-context learning](#) theory from [Section 06.7](#) explains why few-shot prompting works at a mechanistic level.

Prerequisites

This section assumes familiarity with the chat completions API format from [Section 10.1](#), particularly the distinction between system, user, and assistant messages. Understanding how LLMs generate text token by token from [Section 05.1](#) will help you appreciate why prompt wording matters so much for output quality.

11.1.1 The Anatomy of a Prompt



A Mad Libs game card showing a prompt template with blanks to fill in, illustrating parameterized prompts

Figure 11.1.1: Prompt templates are the Mad Libs of AI: fill in the blanks with your specific task, and the model completes the story.

Before diving into specific techniques, it helps to understand the structural components that every prompt can contain. A well-designed prompt typically includes some combination of the following elements: an **instruction** that tells the model what to do, **context** that provides background information, **input data** that the model should process, **output format** specification that constrains the response shape, and **examples** that demonstrate the desired behavior. Not every prompt needs all five components, but being intentional about which components to include (and which to omit) is the first step toward reliable results.

In the [chat API format](#) used by most modern LLMs, these components are distributed across different message roles. The `system` message typically carries the instruction, context, and format specification. The `user` message carries the input data. And when using few-shot examples, pairs of `user` and `assistant` messages demonstrate the expected behavior before the actual query arrives. Figure 11.1.1 maps these structural components to their corresponding chat API roles.

Fun Note

Prompt engineering is one of the few engineering disciplines where adding a polite "please" to your input can measurably improve output quality. Researchers have found that phrasing instructions as respectful requests rather than blunt commands produces better results on many benchmarks. Whether this reflects training data bias or something deeper remains an open question, but it does mean that your manners are now a professional skill.




Figure 11.1.2: The structural components of a prompt mapped to chat API message roles.

11.1.2 Zero-Shot Prompting



Different students in a classroom representing different prompting strategies: zero-shot, few-shot, and chain-of-thought

Figure 11.1.3: Prompting strategies are like study methods: zero-shot wings it, few-shot learns from examples, and chain-of-thought works through the problem step by step.

Zero-shot prompting is the simplest approach: you give the model an instruction and input with no examples. The model relies entirely on its pre-trained knowledge to produce the answer. This works surprisingly well for tasks the model has seen during training, such as summarization, translation, classification of common categories, and question answering. The key to effective zero-shot prompting is **specificity**. Vague instructions produce vague results.

11.1.2.1 The Specificity Principle

Consider the difference between a vague prompt and a specific one for the same task:

Code Fragment 11.1.2 contrasts a vague prompt with a constrained one, showing how specificity transforms the same task from ambiguous to precise.

```
# Comparing vague vs. specific zero-shot prompts
# Specificity in instructions directly controls output quality
# Vague zero-shot prompt
vague_prompt = "Summarize this article."
# Specific zero-shot prompt
specific_prompt = ""Summarize the following article in exactly 3 bullet
points.
Each bullet point should be one sentence of at most 20 words.
Focus on the key findings, not background information.
Use present tense throughout.
Article:
{article_text}""
```

► OUTPUT

```
This product exceeded all my expectations! => positive
The delivery was late and the item arrived damaged. => negative
The package arrived on Tuesday as scheduled. => neutral
```

Code Fragment 11.1.1: *Comparing vague vs. specific zero-shot prompts*

The specific prompt constrains the output length (3 bullet points), format (one sentence each), content focus (key findings), and style (present tense). These constraints dramatically reduce the space of possible outputs, making the model far more likely to produce exactly what you need. This principle applies universally: the more precisely you define success, the more reliably the model achieves it.

11.1.2.2 Zero-Shot Classification Example

Code Fragment 11.1.5 illustrates a chat completion call.


```

# Zero-shot sentiment classification via the chat completions API.
# The system prompt defines the task; no examples are provided.
import openai
# Initialize OpenAI client (reads OPENAI_API_KEY from env)
client = openai.OpenAI()
def classify_sentiment(text: str) -> str:
# Send chat completion request to the API
response = client.chat.completions.create(
model="gpt-4o-mini",
messages=[
{"role": "system",
"content": """"Classify the sentiment of the given text.
Respond with exactly one word: positive, negative, or neutral.
Do not include any explanation or punctuation."""},
{"role": "user",
"content": text}
],
temperature=0.0,
max_tokens=5
)
# Extract the generated message from the API response
return response.choices[0].message.content.strip().lower()
# Test examples
texts = [
"This product exceeded all my expectations!",
"The delivery was late and the item arrived damaged.",
"The package arrived on Tuesday as scheduled."
]
for t in texts:
print(f"{t[:50]:50s} => {classify_sentiment(t)}")

```

Code Fragment 11.1.2: *Zero-shot sentiment classification via the chat completions API.*

Note: Temperature for Classification

Setting `temperature=0.0` for classification tasks ensures deterministic outputs. Since we want a single correct label (not creative variation), zero temperature eliminates randomness in sampling. Combined with `max_tokens=5`, this constrains the model to output only the label.

11.1.3 Few-Shot Prompting



A chef analogy comparing zero-shot (cooking without a recipe) to few-shot (cooking with example dishes to reference)

Figure 11.1.4: *Zero-shot is cooking without a recipe; few-shot is cooking with a few example dishes to reference. More examples usually means tastier results.*

Fun Fact

Few-shot prompting is essentially "teaching by example" compressed into a few sentences. The [GPT-3](#) paper showed that providing just two or three examples can replace hours of [fine-tuning](#). It turns out that LLMs are excellent pattern matchers: show them what you want twice, and they will extrapolate the rest. If only human interns worked the same way.

Key Insight

Why few-shot prompting works (in-context learning): Few-shot examples do not update the model's weights. Instead, they exploit a phenomenon called **in-context learning**: the model's [attention mechanism](#) identifies the pattern in the examples and applies it to the new query, all within a single forward pass. Mechanistically, the transformer's attention heads "soft-retrieve" relevant information from the examples to condition the generation of the answer. This is why example quality matters so much: the model is pattern-matching on whatever structure it finds in the examples, including formatting, reasoning style, and label assignments. Research covered in [Section 06.7](#) explores the theoretical foundations of in-context learning.

Few-shot prompting provides examples of the desired input-output mapping before the actual query. This technique was popularized by the GPT-3 paper (Brown et al., 2020), which showed that providing as few as two or three examples can dramatically improve performance on tasks the model has not been explicitly [fine-tuned](#) for. Few-shot examples serve multiple purposes: they demonstrate the expected output format, clarify

ambiguous instructions, establish the decision boundary for classification tasks, and prime the model's internal representations toward the target behavior.

11.1.3.1 Designing Effective Few-Shot Examples

Not all examples are equally useful. Research and practice have identified several principles for selecting few-shot examples:

- **Cover the label space:** Include at least one example per category. If you have three sentiment labels, show at least one positive, one negative, and one neutral example.
- **Include edge cases:** Show examples near the decision boundary. A mildly negative review is more informative than an obviously negative one.
- **Match the distribution:** If 80% of your real inputs are neutral, your examples should reflect that proportion (or slightly oversample rare classes).
- **Keep formatting consistent:** Every example should follow exactly the same structure. Inconsistency in formatting confuses the model.
- **Order matters:** Recent research shows that example order affects performance. Placing the most relevant example last (closest to the query) often helps.

11.1.3.2 Few-Shot Entity Extraction

Code Fragment 11.1.2 illustrates a chat completion call.

```

# Few-shot entity extraction with edge-case examples
# Demonstrates how to include examples with empty result fields
import openai, json
client = openai.OpenAI()
def extract_entities(text: str) -> dict:
    response = client.chat.completions.create(
        model="gpt-4o-mini",
        messages=[
            {"role": "system",
             "content": """"Extract named entities from the text.
             Return a JSON object with keys: persons, organizations, locations.
             Each value is a list of strings. If no entities are found for a
             category, return an empty list.""""},
            # Few-shot example 1
            {"role": "user",
             "content": "Tim Cook announced the new iPhone at Apple Park in Cupertino."},
            {"role": "assistant",
             "content": '{"persons": ["Tim Cook"], "organizations": ["Apple"], "locations": ["Apple Park", "Cupertino"]}'},
            # Few-shot example 2 (edge case: no persons)
            {"role": "user",
             "content": "The European Central Bank raised interest rates in Frankfurt."},
            {"role": "assistant",
             "content": '{"persons": [], "organizations": ["European Central Bank"], "locations": ["Frankfurt"]}'},
            # Actual query
            {"role": "user",
             "content": text}
        ],
        temperature=0.0
    )
    return json.loads(response.choices[0].message.content)
result = extract_entities(
    "Satya Nadella confirmed that Microsoft will open a new office in Tokyo."
)
print(json.dumps(result, indent=2))

```

► OUTPUT

```

{
  "persons": ["Satya Nadella"],
  "organizations": ["Microsoft"],
  "locations": ["Tokyo"]
}

```

Code Fragment 11.1.3: *Few-shot entity extraction with edge-case examples*

Library Shortcut: Instructor for Guaranteed Structured Output

The few-shot approach above requires manual JSON parsing and careful example engineering to prevent malformed output. The `instructor` library patches your OpenAI client to return validated Pydantic models directly, eliminating JSON parsing errors entirely:

```
# Instructor: extract structured entities directly as a Pydantic model.
# Replaces manual JSON parsing with type-safe, validated output.
import instructor
from openai import OpenAI
from pydantic import BaseModel

class Entities(BaseModel):
    persons: list[str]
    organizations: list[str]
    locations: list[str]

client = instructor.from_openai(OpenAI())
entities = client.chat.completions.create(
    model="gpt-4o-mini",
    response_model=Entities,
    messages=[{"role": "user", "content":
        "Satya Nadella confirmed that Microsoft will open a new office in
        Tokyo."}]
)
print(entities.persons) # ['Satya Nadella']
print(entities.organizations) # ['Microsoft']
print(entities.locations) # ['Tokyo']
```

► OUTPUT

You are a customer support intent classifier.
Classify the customer message into one of these categories: billing,
technical_support, account, general_inquiry
Respond with only the category name, nothing else.

I can't log into my account and I need to reset my password

Code Fragment 11.1.4: *Instructor: extract structured entities directly as a Pydantic model.*

```
pip install instructor
```

 . Instructor handles retries on validation failure, supports nested models, and works with OpenAI, Anthropic, Gemini, and local models via LiteLLM. No few-shot examples needed for this task; the Pydantic schema tells the model exactly what to produce. See [Section 10.2](#) for more on structured output patterns.

Notice how the second few-shot example deliberately shows an edge case where no persons are present. This teaches the model to return an empty list rather than hallucinating a name. Without this example, models sometimes invent a person associated with the European Central Bank.

11.1.4 System Prompts and Role Assignment



A film director (system prompt) giving instructions to an actor (model) about how to play their role

Figure 11.1.5: *The system prompt is the director of the show: it sets the scene, defines the character, and tells the model exactly how to perform.*

The system prompt is the most powerful lever in prompt design. It sets the model's persona, establishes behavioral constraints, defines the output format, and provides persistent context that applies to every turn of the conversation. A well-crafted system prompt transforms a general-purpose model into a specialized tool.

11.1.4.1 Role Prompting

Assigning a role to the model activates domain-specific knowledge and adjusts the style, vocabulary, and reasoning patterns of the response. When you tell a model to act as a "senior data scientist," it tends to provide more technically rigorous answers with appropriate caveats. When you assign the role of "elementary school teacher," it simplifies language and uses analogies. This works because the model has seen text written by people in these roles during [pre-training](#), and the role assignment shifts the probability distribution toward that style of text.

Key Insight

Role prompting is not magic. The model does not actually become an expert. Instead, it shifts its output distribution toward the kind of text that a person in that role would write. This means role prompting works best for roles that are well-represented in the training data (e.g., "Python developer," "medical doctor") and less well for niche or fictional roles. Always validate the output against ground truth rather than trusting the persona.

11.1.4.2 System Prompt Architecture

Production system prompts follow a layered structure. Each layer serves a distinct purpose, and the order matters because models pay more attention to instructions that appear early in the prompt and those that appear at the very end. Keep in mind that every prompt consumes tokens from the model's context window (recall the context window concept from [Chapter 04](#)). A system prompt of 2,000 tokens leaves less room for user input, few-shot examples, and the model's response. For models with 4K or 8K context windows, long system prompts can significantly constrain the available space; models with 128K+ windows offer much more headroom, but token cost still scales with prompt length. Code Fragment 11.1.5 demonstrates this layered structure for a medical coding application.


```

# Layered system prompt architecture for production use
# Sections: Role, Task, Constraints, Output Format, Examples
SYSTEM_PROMPT = """
## Role
You are a medical coding assistant specializing in ICD-10 classification.
You have 15 years of experience in health information management.

## Task
Given a clinical note, extract the primary diagnosis and assign the
correct ICD-10 code. If the note is ambiguous, list the top 3 most
likely codes with confidence scores.

## Constraints
- Only use ICD-10-CM codes (not ICD-10-PCS procedure codes)
- Never fabricate codes; if unsure, say "requires manual review"
- Do not provide medical advice or treatment recommendations
- Flag any note that mentions patient safety concerns

## Output Format
Return a JSON object:
{
  "primary_diagnosis": "description",
  "icd10_code": "X00.0",
  "confidence": 0.95,
  "alternatives": [{"code": "X00.1", "confidence": 0.80}],
  "flags": ["safety_concern"] or []
}

## Examples
[Include 2-3 few-shot examples here]
"""

```

Code Fragment 11.1.5: Layered system prompt architecture for a medical coding assistant. The prompt is organized into five distinct sections (Role, Task, Constraints, Output Format, Examples) that enforce ICD-10 classification rules, prevent code fabrication, and specify a structured JSON response schema with confidence scores and safety flags.

Figure 11.1.3 shows how these layers stack together in a production system prompt architecture.




Figure 11.1.6: A production system prompt follows a layered architecture. Models attend most strongly to early and late content.

Why prompt architecture matters for agents. The layered system prompt structure introduced here becomes critical when building [AI agents](#) (Chapter 22). In agent systems, the system prompt must define not only the persona and output format, but also the agent's available tools, decision-making constraints, and safety boundaries. A poorly structured agent prompt leads to tool misuse, hallucinated actions, and safety violations. The layered architecture (role, task, constraints, format, examples) provides a template that scales from simple chatbots to complex multi-tool agents.

11.1.5 Prompt Templates and Variable Injection

In production applications, prompts are rarely static strings. They are templates with placeholders that get filled at runtime with user input, [retrieved context](#), configuration parameters, and dynamic metadata. A prompt template separates the static instruction logic from the dynamic data, making prompts reusable, testable, and version-controllable.

11.1.5.1 Building a Prompt Template System

Code Fragment 11.1.6 defines a prompt template.

```

# Prompt template system: separate static logic from dynamic data
# Enables version control, testing, and runtime parameterization
from string import Template
from dataclasses import dataclass
from typing import Optional
@dataclass
class PromptTemplate:
    """A reusable prompt template with variable injection."""
    name: str
    version: str
    system_template: str
    user_template: str
    def render(self, **kwargs) -> list[dict]:
        """Render the template with provided variables."""
        return [
            {"role": "system",
            "content": self.system_template.format(**kwargs)},
            {"role": "user",
            "content": self.user_template.format(**kwargs)}
        ]
# Define a reusable classification template
classifier = PromptTemplate(
    name="intent_classifier",
    version="1.2.0",
    system_template="""You are a customer support intent classifier.
Classify the customer message into one of these categories: {categories}
Respond with only the category name, nothing else.""",
    user_template="{message}"
)
# Render at runtime
messages = classifier.render(
    categories="billing, technical_support, account, general_inquiry",
    message="I can't log into my account and I need to reset my password"
)
print(messages[0]["content"])
print(messages[1]["content"])

```

Code Fragment 11.1.6: *Prompt template system: separate static logic from dynamic data*

Warning: Template Injection

When injecting user-provided variables into prompt templates, be aware of prompt injection risks. A malicious user could provide input like "Ignore all previous instructions and..." as their message. [Section 11.4](#) covers defense strategies in detail. As a basic precaution, always validate and sanitize user inputs before template injection, and consider wrapping user content in delimiters like XML tags (`<user_input>...</user_input>`) to help the model distinguish instructions from data.

11.1.6 Handling Edge Cases

Even well-designed prompts encounter edge cases. Understanding the common failure modes and building defenses into your prompts is essential for production reliability.

11.1.6.1 Hallucinations

Models sometimes generate plausible-sounding but factually incorrect information. To mitigate this, explicitly instruct the model to acknowledge uncertainty:

- Add instructions like: "If you are not certain, say 'I don't have enough information to answer this.'"
- Request citations or sources, then verify them independently.
- Use constrained output formats (like selection from a fixed list) to prevent open-ended fabrication.
- Lower the temperature to reduce creative (and potentially inaccurate) responses.

11.1.6.2 Refusals

Models sometimes refuse to answer legitimate queries because they misidentify them as harmful. When building prompts for applications where false refusals are costly, include explicit permission statements in the system prompt. For example: "You are a medical education tool. You may discuss medical conditions, symptoms, and treatments in an

educational context." This helps calibrate the model's safety filters without disabling them entirely.

11.1.6.3 Verbosity Control

Without explicit length constraints, models tend to over-explain. Several techniques help control output length:

- **Word/sentence limits:** "Answer in at most 2 sentences."
- **Format constraints:** "Respond with only the JSON object, no explanation."
- **Negative instructions:** "Do not include any caveats, disclaimers, or preamble."
- **max_tokens parameter:** Set a hard cap at the API level as a safety net.

11.1.7 Iterative Prompt Refinement: A Practical Workflow

Prompt engineering is inherently iterative. The most effective workflow follows a build-measure-improve cycle. Start with a simple prompt, evaluate it against a test set, identify failure patterns, refine the prompt to address those failures, and re-evaluate. Each iteration should change only one aspect of the prompt so you can attribute improvements (or regressions) to specific modifications.

Iteration Comparison

ITERATION	CHANGE	ACCURACY	OBSERVATION
v1	Basic zero-shot: "Classify sentiment"	72%	Many neutral texts misclassified as positive
v2	Added output constraint: "positive, negative, or neutral"	81%	Fewer hallucinated labels, still struggles with sarcasm
v3	Added 3 few-shot examples (including sarcastic review)	89%	Sarcasm handling improved; some edge cases with mixed sentiment
v4	Added instruction: "Focus on the overall tone, not individual phrases"	93%	Mixed-sentiment cases resolved; close to human baseline

Key Insight

This iterative table is not hypothetical. The pattern of 15-20 percentage point improvements from basic zero-shot to well-engineered prompts is consistently observed in practice. The lesson is clear: prompt quality is not binary. It exists on a spectrum, and systematic refinement delivers measurable gains at each step.

Self-Check

Q1: What is the primary advantage of few-shot prompting over zero-shot prompting?

► [Show Answer](#)

Q2: Why is setting temperature=0.0 recommended for classification tasks?

► [Show Answer](#)

Q3: In the system prompt architecture, why does the order of layers matter?

► [Show Answer](#)

Q4: A zero-shot classifier achieves 72% accuracy. After adding few-shot examples and output constraints, it reaches 93%. What should you try next?

► [Show Answer](#)

Q5: What is the risk of injecting user-provided text directly into a prompt template without sanitization?

► [Show Answer](#)

Tip: Put Instructions Before Context

Place your task instructions at the beginning of the prompt, before any long context or documents. Models attend most strongly to the start and end of the prompt. Burying instructions in the middle of a long context often causes them to be ignored.

Key Takeaways

- **Specificity is the foundation of prompt quality.** Vague instructions produce vague outputs. Constrain the format, length, style, and content focus explicitly.
- **Few-shot examples** are the most reliable way to improve accuracy. Include examples that cover the label space, demonstrate edge cases, and maintain consistent formatting.
- **System prompts should follow a layered architecture:** role, task, constraints, output format, then examples. This structure is reusable across applications.
- **Prompt templates separate logic from data.** Use parameterized templates for production applications to enable [version control](#), testing, and reuse.
- **Prompt engineering is iterative.** Start simple, measure against a test set, identify failure patterns, refine one element at a time, and re-measure. Each cycle yields measurable improvements.
- **Build defenses into your prompts from the start.** Handle hallucinations with uncertainty acknowledgment, control verbosity with explicit constraints, and protect against injection with input sanitization.

Real-World Scenario: Redesigning Prompts to Fix a Content Moderation Pipeline

Who: A trust and safety engineer at a social media startup using GPT-4 to classify user-generated posts as safe, borderline, or unsafe.

Situation: The initial prompt was a single sentence: "Classify this post as safe, borderline, or unsafe." The system processed 80,000 posts per day with results fed into an automated moderation queue.

Problem: The classifier had a 23% disagreement rate with human moderators, with most errors being false negatives on subtle toxicity (sarcasm, coded language) and false positives on discussions about sensitive topics (mental health support, news about violence).

Dilemma: They could fine-tune a model on their moderation data (expensive, slow iteration), add more few-shot examples to the prompt (quick but limited by context window), or restructure the entire prompt using a layered architecture with explicit criteria and edge-case examples.

Decision: They restructured the prompt using a layered system prompt: role definition ("You are a content moderation specialist"), explicit criteria for each category with boundary definitions, five few-shot examples covering edge cases (sarcasm, news reporting, support language), and a required confidence score.

How: They built a prompt template with placeholders for the post content, used the system prompt for the role and criteria layers, and placed few-shot examples in the first user/assistant turns. They iterated over three rounds, each time analyzing the 50 highest-disagreement examples from the previous batch.

Result: Disagreement with human moderators dropped from 23% to 7% over three iteration cycles spanning two weeks. False negatives on coded toxicity fell by 60%, and false positives on sensitive-but-safe content fell by 75%.

Lesson: A structured, layered prompt with explicit criteria and edge-case examples consistently outperforms vague instructions; iterating against your specific failure patterns is more effective than generic prompt improvements.

Research Frontier

Automatic prompt optimization. Tools like DSPy (Khattab et al., 2024) and OPRO (Yang et al., 2024) treat prompt engineering as an optimization problem, automatically searching the space of prompt variations to maximize task performance. This shifts prompt design from manual craft toward systematic search.

Prompt sensitivity research. Studies show that semantically equivalent prompts can produce wildly different accuracy scores (sometimes varying by 20% or more), and that optimal prompt format varies across models and model sizes. This motivates systematic prompt evaluation rather than relying on intuition alone.

Cross-model prompt transfer. Research into whether prompts optimized for one model transfer effectively to others shows mixed results: simple prompts transfer well, but complex few-shot examples and formatting choices often need re-optimization per model family.

Exercises

Exercise 11.1.1: Prompt anatomy

Name the five structural components of a well-designed prompt and explain which chat API message role typically carries each component.

► Answer Sketch

Exercise 11.1.2: Zero-shot vs. few-shot

Write two versions of a sentiment classification prompt: one zero-shot and one three-shot. Both should classify customer reviews as positive, neutral, or negative and return JSON.

► [Answer Sketch](#)

Exercise 11.1.3: Role prompting effectiveness

A developer writes: 'You are a helpful assistant. Answer the following medical question.' Explain two problems with this role assignment and rewrite it to be more effective.

► [Answer Sketch](#)

Exercise 11.1.4: System prompt design

Design a system prompt for a customer service chatbot that handles order status inquiries. Include role assignment, behavioral constraints, output format, and escalation rules. Keep it under 200 tokens.

► [Answer Sketch](#)

Exercise 11.1.5: Template parameterization

You have a prompt template with 5 variables that works well for English product reviews. When you switch to French reviews, accuracy drops from 92% to 74%. Identify two likely causes and propose a fix for each.

► [Answer Sketch](#)



References & Further Reading

FOUNDATIONAL RESEARCH

[Brown, T. et al. \(2020\). Language Models are Few-Shot Learners. NeurIPS 2020.](#)

The GPT-3 paper that demonstrated in-context learning at scale. Introduced the few-shot paradigm where examples in the prompt guide model behavior without gradient updates. Essential context for understanding why prompt design matters.

 PAPER

[Min, S. et al. \(2022\). Rethinking the Role of Demonstrations: What Makes In-Context Learning Work? EMNLP 2022.](#)

Surprising finding that the label correctness of few-shot examples matters less than their format and input distribution. Changes how practitioners think about example selection and challenges the intuition that examples must be perfectly labeled.

 PAPER

[Liu, J. et al. \(2022\). What Makes Good In-Context Examples for GPT-3? DeeLIO Workshop, ACL 2022.](#)

Investigates how example selection and ordering affect few-shot performance. Proposes retrieval-based example selection using semantic similarity, showing significant gains over random selection.

 PAPER

PRACTICAL GUIDES

[OpenAI. \(2024\). Prompt Engineering Guide.](#)

OpenAI's official prompt engineering guide with practical strategies for writing clear instructions, providing reference text, and splitting complex tasks. Regularly updated with new model capabilities.

 TUTORIAL

[Anthropic. \(2024\). Prompt Engineering Documentation.](#)

Anthropic's guide to prompting Claude, covering system prompts, XML tags for structure, and techniques specific to Claude's training. Useful companion to the OpenAI guide for

cross-provider prompt design.

 TUTORIAL

[Zamfirescu-Pereira, J. D. et al. \(2023\). Why Johnny Can't Prompt: How Non-AI Experts Try \(and Fail\) to Design LLM Prompts. CHI 2023.](#)

HCI study revealing common failure patterns when non-experts write prompts, including vague instructions, missing constraints, and overreliance on single attempts. Valuable for understanding why systematic prompt design is necessary.

 PAPER

Chain-of-Thought & Reasoning Techniques

Teaching models to think step by step before answering

Show your work. It is not just good advice for students; it turns out to be good advice for language models too.



Prompt

, Show-Your-Work AI Agent

Big Picture

Why reasoning techniques matter. Standard prompting asks a model to jump directly from question to answer. For simple factual lookups, this works fine. But for multi-step reasoning, math problems, logic puzzles, and complex analysis, direct answering leads to frequent errors. Chain-of-Thought (CoT) prompting, introduced by Wei et al. (2022), showed that simply asking the model to "think step by step" before answering can dramatically improve accuracy on reasoning tasks. Building on the foundational techniques from [Section 11.1](#), this section covers CoT and its successors: self-consistency (sample multiple reasoning paths and vote), Tree-of-Thought (structured exploration with backtracking), step-back prompting, and the [ReAct](#) framework that interleaves reasoning with [tool use](#).

Prerequisites

This section builds on the foundational prompt patterns from [Section 11.1](#) (zero-shot, few-shot, role prompting). Understanding of how [attention mechanisms](#) process context from [Section 03.2](#) and the decoding strategies from [Section 05.1](#) will help explain why chain-of-thought prompting improves reasoning quality.



A branching tree structure showing multiple reasoning paths being explored and evaluated simultaneously

Figure 11.2.1: *Tree-of-Thought: instead of one linear chain, the model explores multiple reasoning branches and picks the most promising path.*

11.2.1 Chain-of-Thought Prompting



A student showing their work on a math test, analogous to chain-of-thought prompting where the model explains its reasoning steps

Figure 11.2.2: *Chain-of-thought prompting: making the model show its work, just like your math teacher always insisted you do.*

Paper Spotlight: Wei et al. (2022)

"Chain-of-Thought Prompting Elicits Reasoning in Large Language Models" demonstrated that providing step-by-step reasoning examples in the prompt dramatically improves performance on arithmetic, commonsense, and symbolic reasoning tasks. The key contribution was showing that reasoning ability is an emergent property: it appears in models above ~100B parameters and is absent in smaller models. This paper launched an entire subfield of reasoning-oriented prompt engineering.

Chain-of-Thought prompting works by encouraging the model to generate intermediate reasoning steps before producing a final answer. The mechanism is straightforward: each reasoning step conditions the generation of subsequent steps, allowing the model to "carry" information forward through the computation. Without CoT, the model must compress all reasoning into a single forward pass through the network. With CoT, the model effectively uses its own generated text as a scratchpad, offloading intermediate computation into the token sequence.

Fun Note

The original "Let's think step by step" prompt that launched chain-of-thought research is exactly six words long. Those six words improved GSM8K math accuracy from 18% to 57% on PaLM 540B. Per character, it may be the highest-ROI engineering intervention in the history of AI. Prompt engineers everywhere took note: sometimes the best optimization is just asking nicely.

Aha Moment: CoT as Working Memory

Without CoT, the model must fit all reasoning into a single forward pass through the [transformer](#). With CoT, the model's own output tokens become **working memory**. Each generated reasoning step is literally fed back into the model as input for the next step, just as you would use a scratchpad to carry intermediate results. This is why CoT helps on multi-step problems: it gives the model a place to store intermediate results that its fixed-size hidden state cannot hold all at once.

11.2.1.1 Zero-Shot CoT

The simplest form of CoT requires no examples at all. Kojima et al. (2022) discovered that appending the phrase "Let's think step by step" to a prompt is sufficient to trigger reasoning behavior in large models. (For models with built-in reasoning, such as o3 and DeepSeek-R1, see [Section 11.5](#) on prompting [reasoning models](#).) This zero-shot CoT approach is remarkably effective, improving accuracy on GSM8K (grade school math) from 17.7% to 78.7% for PaLM 540B.

Code Fragment 11.2.2 illustrates a chat completion call.


```
# Zero-shot Chain-of-Thought: append "step by step" to trigger reasoning.
# The model shows its work before producing a final ANSWER: line.
import openai
# Initialize OpenAI client (reads OPENAI_API_KEY from env)
client = openai.OpenAI()
def solve_with_cot(problem: str) -> str:
    """Solve a problem using zero-shot Chain-of-Thought."""
    response = client.chat.completions.create(
        model="gpt-4o",
        messages=[
            {"role": "system",
             "content": """Solve the problem step by step.
Show your reasoning clearly, then provide the final answer
on a separate line starting with "ANSWER: ". """},
            {"role": "user",
             "content": problem}
        ],
        temperature=0.0
    )
    # Extract the generated message from the API response
    return response.choices[0].message.content
problem = """A store sells notebooks for $4 each. If you buy 5 or more,
you get a 20% discount on the entire purchase. Sarah buys 7 notebooks
and pays with a $50 bill. How much change does she receive?"""
print(solve_with_cot(problem))
```

► OUTPUT

Let me work through this step by step.

Step 1: Calculate the base price.

7 notebooks at \$4 each = $7 \times \$4 = \28

Step 2: Check if the discount applies.

Sarah is buying 7 notebooks, which is 5 or more, so the 20% discount applies.

Step 3: Calculate the discount amount.

20% of \$28 = $0.20 \times \$28 = \5.60

Step 4: Calculate the discounted total.

$\$28 - \$5.60 = \$22.40$

Step 5: Calculate the change from \$50.

$\$50 - \$22.40 = \$27.60$

ANSWER: \$27.60

Code Fragment 11.2.1: *Zero-shot Chain-of-Thought: append "step by step" to trigger reasoning.*

11.2.1.2 Few-Shot CoT

For tasks where zero-shot CoT is insufficient, providing examples of step-by-step reasoning significantly improves performance. The key insight is that the examples teach the model not just the answer format, but the *style of reasoning* to apply. Different reasoning styles suit different tasks: arithmetic problems benefit from sequential calculation steps, logic problems benefit from listing premises and drawing inferences, and coding problems benefit from planning before implementing.

Code Fragment 11.2.2 illustrates a chat completion call.

```

# Few-shot Chain-of-Thought: exemplar reasoning chains teach the model
# the style of step-by-step problem solving to apply.
import openai
client = openai.OpenAI()
# Few-shot CoT: provide exemplar reasoning chains
FEW_SHOT_EXAMPLES = """
Example 1:
Q: A train travels 120 miles in 2 hours. It then travels 90 miles in 1.5 hours.
What is the average speed for the entire trip?
A: Let me work through this step by step.
Total distance = 120 + 90 = 210 miles
Total time = 2 + 1.5 = 3.5 hours
Average speed = total distance / total time = 210 / 3.5 = 60 mph
ANSWER: 60 mph
Example 2:
Q: A store has a "buy 2, get 1 free" deal on $6 shirts. How much do 7 shirts
cost?
A: Let me work through this step by step.
For every 3 shirts, you pay for 2: that is 2 x $6 = $12 per group of 3.
7 shirts = 2 complete groups (6 shirts) + 1 remaining shirt.
Cost for 6 shirts = 2 x $12 = $24
Cost for 1 extra shirt = $6
Total = $24 + $6 = $30
ANSWER: $30
"""

def solve_with_few_shot_cot(problem: str) -> str:
    # Send chat completion request to the API
    response = client.chat.completions.create(
        model="gpt-4o",
        messages=[
            {"role": "system",
             "content": f"""Solve problems step by step, following the style shown in these
examples:
{FEW_SHOT_EXAMPLES}
Show your reasoning clearly, then provide the final answer on a line starting
with "ANSWER: ".""""},
            {"role": "user", "content": problem}
        ],
        temperature=0.0
    )
    # Extract the generated message from the API response
    return response.choices[0].message.content
problem = """A bakery sells cupcakes for $3 each. If you buy a dozen (12),

```

```
you get a 25% discount. Maria buys 15 cupcakes. How much does she pay?"""
print(solve_with_few_shot_cot(problem))
```

► OUTPUT

Let me work through this step by step.
First, I need to figure out the pricing structure.
Individual price: \$3 per cupcake.
Dozen deal: 12 cupcakes at 25% off.

Maria buys 15 cupcakes. She can buy 1 dozen (12) at the discount, plus 3 at full price.

Cost for the dozen: $12 \times \$3 = \36 , then 25% off: $\$36 \times 0.75 = \27
Cost for 3 extra: $3 \times \$3 = \9
Total: $\$27 + \$9 = \$36$

ANSWER: \$36

Code Fragment 11.2.2: *Few-shot Chain-of-Thought: exemplar reasoning chains teach the model*

Key Insight

CoT prompting improves accuracy primarily on tasks that require multi-step reasoning. For single-step tasks (simple factual recall, sentiment classification), CoT adds unnecessary tokens without improving quality and may even slightly reduce accuracy due to the additional opportunity for the model to "talk itself into" a wrong answer. The rule of thumb: if a human needs a scratchpad to solve the problem, CoT will help. If a human can answer instantly, CoT is unnecessary overhead.

Figure 11.2.1 contrasts direct prompting with chain-of-thought, showing how decomposition into intermediate steps enables verification at each stage.




Figure 11.2.3: Direct prompting compresses all reasoning into one step; CoT decomposes it into verifiable intermediate steps.

Why does chain-of-thought work mechanistically? The key insight is that transformers compute each output token conditioned on all previous tokens. When the model generates reasoning steps as intermediate tokens, those tokens become part of the input for subsequent generation. In effect, the model is augmenting its own context window with intermediate computation results. Without CoT, a model tackling " $17 \times 23 + 45$ " must compute the entire answer in a single forward pass through its layers. With CoT, it first generates " $17 \times 23 = 391$ " as text tokens, and then those tokens are fed back into the model as context for computing " $391 + 45 = 436$." Each forward pass handles a simpler sub-problem. This is why CoT helps specifically on multi-step problems: it decomposes a hard problem into a sequence of easy problems, where each step's output becomes the next step's input.

11.2.2 Self-Consistency



A jury of multiple model outputs voting on the best answer, illustrating self-consistency sampling

Figure 11.2.4: Self-consistency assembles a jury of model outputs and lets them vote. The majority answer is usually the most reliable one.

Fun Fact

Self-consistency works by asking the model the same question multiple times and taking a majority vote. It is the "ask the audience" lifeline from Who Wants to Be a Millionaire, except every audience member is the same model running at different temperatures. Surprisingly, this simple trick boosted GSM8K math accuracy by over 10 percentage points.

Paper Spotlight: Wang et al. (2022)

"Self-Consistency Improves Chain of Thought Reasoning in Language Models" showed that sampling multiple reasoning paths and taking a majority vote dramatically outperforms greedy single-path CoT. On GSM8K, self-consistency pushed accuracy from 78.7% to over 90% for PaLM 540B. The key insight: correct reasoning paths tend to converge on the same answer, while errors are randomly distributed across wrong answers.

Self-consistency, introduced by Wang et al. (2022), builds on CoT by [sampling](#) multiple reasoning paths at non-zero [temperature](#) and selecting the answer that appears most frequently. The intuition is that while any single reasoning chain might contain errors, different chains are likely to make different errors. When multiple independent chains converge on the same answer, confidence in that answer increases.

11.2.2.1 Implementation

Code Fragment 11.2.6 illustrates a chat completion call.

```

# Self-consistency: sample multiple CoT paths, then majority-vote
# Higher temperature (0.7) ensures diverse reasoning paths
import openai
from collections import Counter
# Initialize OpenAI client (reads OPENAI_API_KEY from env)
client = openai.OpenAI()
def solve_with_self_consistency(problem: str, n_samples: int = 5) -> dict:
    """Sample multiple CoT paths and take majority vote."""
    answers = []
    for _ in range(n_samples):
        # Send chat completion request to the API
        response = client.chat.completions.create(
            model="gpt-4o",
            messages=[
                {"role": "system",
                 "content": """Solve step by step. End with ANSWER: <number>"""},
                {"role": "user", "content": problem}
            ],
            temperature=0.7 # Higher temperature for diverse paths
        )
        # Extract the generated message from the API response
        text = response.choices[0].message.content
        # Extract the final answer
        if "ANSWER:" in text:
            answer = text.split("ANSWER:")[-1].strip()
            answers.append(answer)
        # Majority vote
        vote_counts = Counter(answers)
        best_answer, count = vote_counts.most_common(1)[0]
        return {
            "answer": best_answer,
            "confidence": count / len(answers),
            "all_answers": dict(vote_counts)
        }
result = solve_with_self_consistency(
    "If a train travels at 60 mph for 2.5 hours, then at 80 mph for 1.5 hours, "
    "what is the total distance traveled?"
)
print(f"Answer: {result['answer']}")
print(f"Confidence: {result['confidence']:.0%}")
print(f"All votes: {result['all_answers']}")

```

► OUTPUT

```
Answer: 270 miles
Confidence: 100%
All votes: {'270 miles': 5}
```

Code Fragment 11.2.3: *Self-consistency: sample multiple CoT paths, then majority-vote*

Note: Temperature and Self-Consistency

Self-consistency requires `temperature > 0` to generate diverse reasoning paths. If temperature is zero, every sample produces identical output, defeating the purpose. A temperature of 0.5 to 0.7 provides a good balance between diversity and quality. Higher temperatures produce more diverse paths but increase the chance of individually nonsensical reasoning chains.

11.2.3 Tree-of-Thought (ToT)

Tree-of-Thought (Yao et al., 2023) extends CoT by exploring multiple reasoning paths in a structured tree, with the ability to evaluate and backtrack. While CoT follows a single linear chain and self-consistency samples independent chains, ToT builds a branching exploration tree where the model can evaluate partial solutions, prune unpromising branches, and explore alternatives.

Misconception Warning: ToT Is Not Practical for Most Production Use

Tree-of-Thought is an elegant research technique, but it requires **10 to 50+ API calls per query** (one for each node explored in the tree). At \$0.01 per call, a single ToT query can cost \$0.10 to \$0.50. For production workloads, self-consistency (3 to 5 calls) achieves most of the accuracy benefit at a fraction of the cost. Use ToT only for high-value, low-volume tasks like planning, puzzle solving, or complex [code generation](#) where the cost per query is justified.

The ToT framework has three core components:

1. **Thought generation:** At each step, the model proposes multiple possible next steps (branches).
2. **Thought evaluation:** The model (or a separate evaluator) scores each branch on how promising it looks.
3. **Search strategy:** A search algorithm (typically breadth-first or depth-first) navigates the tree, expanding the most promising nodes and [pruning](#) dead ends.

Figure 11.2.2 illustrates this branching exploration, where each node is evaluated and only promising paths are expanded further.




Figure 11.2.5: *Tree-of-Thought generates multiple branches at each step, evaluates them, prunes poor candidates, and backtracks when needed.*

11.2.3.1 Simplified ToT Implementation

Code Fragment 11.2.4 illustrates a chat completion call.

```

# Tree-of-Thought: generate candidate steps, evaluate each, prune weak branches
# Uses a separate evaluation call to score reasoning quality at each step
import openai
# Initialize OpenAI client (reads OPENAI_API_KEY from env)
client = openai.OpenAI()
def generate_thoughts(problem: str, context: str, n: int = 3) -> list[str]:
    """Generate n candidate next steps."""
    # Send chat completion request to the API
    response = client.chat.completions.create(
        model="gpt-4o",
        messages=[{"role": "user",
                    "content": f"""Problem: {problem}
Progress so far: {context}
Generate {n} different possible next steps. Number them 1, 2, 3.
Each should be a single reasoning step."""}],
        temperature=0.8
    )
    # Extract the generated message from the API response
    text = response.choices[0].message.content
    return [line.strip() for line in text.split("\n") if line.strip()]
def evaluate_thought(problem: str, thought: str) -> float:
    """Score a thought from 0 to 1."""
    response = client.chat.completions.create(
        model="gpt-4o",
        messages=[{"role": "user",
                    "content": f"""Problem: {problem}
Proposed reasoning step: {thought}
Rate this step from 0.0 to 1.0 based on correctness and
usefulness. Respond with only a number."""}],
        temperature=0.0
    )
    try:
        return float(response.choices[0].message.content.strip())
    except ValueError:
        return 0.5
# Example: solve with tree exploration
problem = "Find two numbers that multiply to 36 and add to 13."
thoughts = generate_thoughts(problem, "(starting)")
for t in thoughts[:3]:
    score = evaluate_thought(problem, t)
    print(f" [{score:.2f}] {t}")

```

► OUTPUT

```
[0.90] 1. List factor pairs of 36: (1,36), (2,18), (3,12), (4,9), (6,6)
[0.70] 2. Set up equations:  $x * y = 36$  and  $x + y = 13$ , solve quadratic
[0.50] 3. Try guess and check starting with numbers close to  $\sqrt{36}$ 
```

Code Fragment 11.2.4: *Tree-of-Thought: generate candidate steps, evaluate each, prune weak branches*

11.2.4 Step-Back Prompting

Step-back prompting (Zheng et al., 2023) takes the opposite approach from diving into details. Instead of reasoning through the specific problem, it first asks the model to identify the abstract principle or high-level concept, then applies that principle to solve the specific problem. This is particularly effective for science, math, and policy questions where the correct reasoning requires recalling a general rule before applying it.

The two-phase approach works as follows. First, generate a "step-back" question: "What general principle or concept is relevant to this problem?" Then, use the model's answer to that abstract question as context for solving the original specific problem. This prevents the model from getting lost in surface-level details and grounds its reasoning in correct foundational knowledge.

Code Fragment 11.2.6 illustrates a chat completion call.

```

# Step-back prompting: abstract the principle first, then solve
# Phase 1 identifies the relevant law; Phase 2 applies it
import openai
# Initialize OpenAI client (reads OPENAI_API_KEY from env)
client = openai.OpenAI()
def step_back_solve(question: str) -> str:
    """Two-phase step-back prompting: abstract first, then solve."""
    # Phase 1: Generate the step-back (abstract) question
    abstraction = client.chat.completions.create(
        model="gpt-4o",
        messages=[
            {"role": "system",
             "content": "Given a specific question, identify the general principle or
             foundational concept needed to answer it. Respond with ONLY the principle, not
             the answer to the original question."},
            {"role": "user", "content": question}
        ],
        temperature=0.0
    )
    principle = abstraction.choices[0].message.content
    print(f"Step-back principle: {principle[:120]}...")
    # Phase 2: Solve using the principle as context
    solution = client.chat.completions.create(
        model="gpt-4o",
        messages=[
            {"role": "system",
             "content": f"Use this foundational principle to answer the
             question:\n\n{principle}"},
            {"role": "user", "content": question}
        ],
        temperature=0.0
    )
    return solution.choices[0].message.content
answer = step_back_solve(
    "If the temperature of an ideal gas is doubled while keeping "
    "the volume constant, what happens to its pressure?"
)
print(f"\nAnswer: {answer}")

```

► OUTPUT

Step-back principle: Gay-Lussac's Law (Pressure-Temperature Law): For an ideal gas at constant volume, pressure is directly...

Answer: According to Gay-Lussac's Law, pressure is directly proportional to absolute temperature when volume is held constant ($P/T = \text{constant}$). If the temperature is doubled (from T to $2T$), the pressure also doubles (from P to $2P$).

Code Fragment 11.2.5: *Step-back prompting: abstract the principle first, then solve*

When to Use Step-Back Prompting

Step-back prompting excels on questions where **domain knowledge recall** is the bottleneck, not multi-step computation. Physics, chemistry, law, and policy questions often benefit because the model needs to recall the correct rule before applying it. For pure arithmetic or logic problems, standard CoT is usually sufficient. Step-back prompting adds one extra LLM call per question, so it doubles latency and cost; use it selectively for high-value queries where accuracy matters more than speed.

11.2.5 The ReAct Framework

Why this matters: ReAct is the bridge between prompt engineering and [agentic AI](#). Up to this point, all prompting techniques operate within the model's existing knowledge. But real-world tasks often require current information (stock prices, weather, database records) that the model cannot know from pre-training. ReAct solves this by interleaving reasoning with tool use: the model thinks about what information it needs, calls a tool to get it, observes the result, and continues reasoning. This pattern is the foundation of the agent architectures covered in [Chapter 22](#).

ReAct (Yao et al., 2022) combines **Reasoning** and **Acting** in an interleaved loop. The model alternates between generating reasoning traces (thinking about what to do) and taking actions (calling tools, searching databases, executing code). This pattern is fundamental to modern LLM agents and represents the bridge between prompt engineering and agentic AI. [Chapter 22](#) explores full agent architectures that build on the ReAct pattern, including multi-turn tool use, planning loops, and error recovery.

11.2.5.1 The ReAct Loop

Code Fragment 11.2.6 illustrates a chat completion call.

```

# ReAct agent: interleave reasoning (THINK) with tool use (ACT)
# Runs up to 5 iterations of search-and-reason before producing a final answer
import openai, json
# Initialize OpenAI client (reads OPENAI_API_KEY from env)
client = openai.OpenAI()
TOOLS = [
{"type": "function",
"function": {
"name": "search",
"description": "Search for information on a topic",
"parameters": {
"type": "object",
"properties": {
"query": {"type": "string", "description": "Search query"}
},
"required": ["query"]
}
}}
]
REACT_SYSTEM = """You are a research assistant. For each question:
1. THINK: Reason about what information you need
2. ACT: Use the search tool to find information
3. OBSERVE: Analyze the search results
4. Repeat THINK/ACT/OBSERVE until you have enough information
5. Provide a final, well-sourced answer"""
def react_agent(question: str) -> str:
messages = [
{"role": "system", "content": REACT_SYSTEM},
{"role": "user", "content": question}
]
for step in range(5): # Max 5 iterations
# Send chat completion request to the API
response = client.chat.completions.create(
model="gpt-4o",
messages=messages,
tools=TOOLS,
tool_choice="auto"
)
# Extract the generated message from the API response
msg = response.choices[0].message
if msg.tool_calls:
# Model wants to use a tool (ACT)
messages.append(msg)
for call in msg.tool_calls:
result = execute_search(

```

```

json.loads(call.function.arguments)["query"]
)
messages.append({
    "role": "tool",
    "tool_call_id": call.id,
    "content": result
})
else:
    # Model has enough info, return final answer
    return msg.content
return "Max iterations reached without final answer."

```

Code Fragment 11.2.6: *ReAct agent loop in `react_agent()` that interleaves reasoning and tool use over up to 5 iterations. The `REACT_SYSTEM` prompt defines the Think/Act/Observe cycle, `tool_choice="auto"` lets the model decide when to search, and the loop terminates when the model produces a final answer without requesting any tool calls.*

Cost Awareness

Advanced reasoning techniques like self-consistency and ToT multiply the number of API calls per query. Self-consistency with 5 samples costs 5x a single CoT call. ToT can require 10 to 50 calls depending on the branching factor and depth. Always consider the cost-accuracy tradeoff. For a batch of 10,000 queries, self-consistency with $n=5$ at \$0.01 per call adds up to \$500 compared to \$100 for single CoT. The improvement from 90% to 95% accuracy must be worth the 5x cost increase for your use case.


11.2.6 Comparison of Reasoning Techniques

🛠️ 6. Comparison of Reasoning Techniques Intermediate

TECHNIQUE	API CALLS	BEST FOR	LIMITATION
Zero-shot CoT	1	Simple multi-step reasoning	Single path can still err
Few-shot CoT	1	Domain-specific reasoning style	Requires good examples
Self-consistency	n (typically 5 to 20)	Math, logic, factual questions	High cost; assumes closed-form answer
Tree-of-Thought	10 to 50+	Complex planning, puzzles	Very expensive; complex to implement
Step-back prompting	2	Science, policy, principle-based questions	Adds latency; not useful for procedural tasks
ReAct	2 to 10+	Questions requiring external information	Requires tool infrastructure

11.2.7 Choosing a Reasoning Technique: Decision Flowchart

With multiple reasoning techniques available, selecting the right one for a given task is itself a decision problem. The following flowchart (Figure 11.2.4) provides a practical starting point. Start at the top and follow the questions to arrive at a recommended technique.



Tip: Consider reasoning models (o3, o4-mini, DeepSeek R1) as alternatives to CoT prompting. These models perform chain-of-thought internally, so explicit CoT prompting may be unnecessary.

Figure 11.2.6: Decision flowchart for selecting a reasoning technique. Start at the top and follow the decision path to find the best technique for your task.

Reasoning Models vs. Prompting for Reasoning

Providers now offer **reasoning models** with built-in chain-of-thought: OpenAI's o3 and o4-mini, DeepSeek R1, and Claude's extended thinking mode. These models perform multi-step reasoning internally without needing explicit CoT prompting. When using a reasoning model, adding "think step by step" to your prompt is redundant and may even degrade quality. The decision is: pay for a more expensive reasoning model that handles reasoning natively, or use a standard model with explicit CoT prompting at lower per-token cost but more prompt engineering effort.

Self-Check

Q1: Why does CoT prompting improve accuracy on math problems but not on simple sentiment classification?

► [Show Answer](#)

Q2: Self-consistency uses temperature > 0 while standard CoT often uses temperature $= 0$. Explain why.

► [Show Answer](#)

Q3: How does Tree-of-Thought differ from running self-consistency multiple times?

► [Show Answer](#)

Q4: In the ReAct framework, what happens if the model enters an infinite loop of searching without producing a final answer?

► [Show Answer](#)

Q5: When would you choose step-back prompting over standard CoT?

► [Show Answer](#)

Modify and Observe

Experiment with the CoT examples from this section:

1. Take the zero-shot CoT example and remove "step by step" from the system prompt. Compare the accuracy on a multi-step math problem. How many problems does it get wrong without the CoT trigger?
2. In the self-consistency example, change the number of samples from 5 to 1, 3, 10, and 20. Plot the accuracy at each sample count. At what point do additional samples stop helping?
3. Try adding CoT to a simple yes/no factual question (e.g., "Is Python a compiled language?"). Observe whether the model "overthinks" and produces a less confident or incorrect answer. This demonstrates why CoT can hurt on simple tasks.

Tip: Use Structured Output Formats

When you need parseable output, explicitly request JSON and provide a schema example in your prompt. Add "Respond ONLY with valid JSON, no other text" to reduce the chance of preamble text. For critical applications, use the API's native JSON mode if available.

Key Takeaways

- **Chain-of-Thought** is the single most impactful prompting technique for reasoning tasks. "Think step by step" can improve math accuracy from under 20% to over 78%, and it costs nothing extra in implementation complexity.
- **Self-consistency trades cost for accuracy.** Sampling 5 to 10 reasoning paths and voting raises accuracy further, but multiplies API cost linearly. Use it when correctness is worth the expense.
- **Tree-of-Thought is powerful but expensive.** Reserve it for complex planning problems where early pruning of bad paths saves overall computation compared to exhaustive self-consistency sampling.
- **Step-back prompting grounds reasoning in principles.** For science, law, and domain-specific reasoning, abstracting before solving prevents the model from getting lost in surface-level details.
- **ReAct bridges prompting and agents.** By interleaving reasoning with tool use, ReAct enables the model to access external information and take actions, forming the foundation of modern AI agent architectures.
- **Always match the technique to the task.** Simple tasks need simple prompts. Complex tasks need complex reasoning strategies. Overengineering wastes cost; underengineering wastes accuracy.

Real-World Scenario: Using Self-Consistency to Improve Medical Triage Accuracy

Who: A health-tech startup's ML team building an AI-assisted symptom checker that suggests urgency levels (emergency, urgent, routine, self-care) for patient-reported symptoms.

Situation: Their chain-of-thought prompt achieved 81% agreement with physician triage decisions on a benchmark of 2,000 cases, but the 19% error rate was unacceptable for a safety-critical application, particularly for false negatives (under-triaging emergencies).

Problem: Single CoT reasoning paths sometimes fixated on the most obvious symptom while missing combinations that indicated higher urgency (e.g., mild chest pain plus shortness of breath plus age over 60).

Dilemma: They could [fine-tune](#) a specialized model (months of work, regulatory implications), add more medical context to the prompt (marginal gains, context window limits), or use self-consistency sampling with majority voting (5x cost increase per query).

Decision: They implemented self-consistency with 7 reasoning paths at temperature 0.7, using weighted majority voting where paths that mentioned more symptoms received slightly higher weight.

How: Each query spawned 7 parallel API calls with the same CoT prompt. A lightweight post-processing step extracted the triage level from each path and applied majority voting. For cases where no category received a majority, the system defaulted to the highest urgency level mentioned.

Result: Agreement with physician decisions rose from 81% to 91%. False negatives (missed emergencies) dropped by 72%, which was the most critical improvement. The 7x cost increase was offset by running the system only for cases where a single-path confidence score fell below 0.85, covering roughly 30% of total queries.

Lesson: Self-consistency is most valuable for high-stakes decisions where a single reasoning path may miss important factors; applying it selectively (only on uncertain cases) controls cost while maximizing safety gains.

Research Frontier

Reasoning tokens and internal chain-of-thought. Models like OpenAI o1/o3 and DeepSeek-R1 internalize chain-of-thought reasoning, generating "thinking" tokens before producing answers. This shifts CoT from a prompting technique to a model capability, as explored in [Section 10.4](#) and [Section 11.5](#).

Self-consistency and verification. Research on self-consistency (Wang et al., 2023) shows that sampling multiple reasoning paths and taking the majority answer significantly improves accuracy, especially on math and logic tasks, at the cost of higher inference compute.

Faithful versus unfaithful reasoning. Studies reveal that CoT explanations do not always reflect the model's actual reasoning process. The model may arrive at the correct answer through shortcuts while generating a plausible-looking explanation. This has implications for using CoT as an [interpretability](#) tool, a topic explored further in [Section 18.1](#).

Exercises

Exercise 11.2.1: CoT mechanism

Explain why Chain-of-Thought prompting improves accuracy on multi-step reasoning problems. What role does the model's own generated text play?

► Answer Sketch

Exercise 11.2.2: Zero-shot CoT implementation

Write a function that takes a math word problem as input and solves it using zero-shot CoT. The function should append 'Let's think step by step' to the prompt and then extract the final numerical answer from the response.

► Answer Sketch

Exercise 11.2.3: Self-consistency

Describe the self-consistency technique. Why does sampling multiple reasoning paths and taking a majority vote improve accuracy compared to a single CoT chain?

► Answer Sketch

Exercise 11.2.4: Tree-of-Thought

Implement a simplified Tree-of-Thought solver for a puzzle problem. Generate 3 candidate partial solutions at each step, evaluate them with a scoring prompt, prune the weakest, and continue with the top 2.

► Answer Sketch

Exercise 11.2.5: ReAct framework

Compare the ReAct pattern (interleaving reasoning and action) with pure Chain-of-Thought. For a question like 'What is the population of the city where the 2024 Olympics were held?', trace through both approaches and identify where pure CoT might fail.

► Answer Sketch



References & Further Reading

CORE CHAIN-OF-THOUGHT PAPERS

[Wei, J. et al. \(2022\). Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. NeurIPS 2022.](#)

The seminal paper that introduced chain-of-thought prompting. Demonstrates that adding "let's think step by step" or providing reasoning examples dramatically improves performance on arithmetic, commonsense, and symbolic reasoning tasks.

 PAPER

[Kojima, T. et al. \(2022\). Large Language Models are Zero-Shot Reasoners. NeurIPS 2022.](#)

Shows that simply appending "Let's think step by step" to a prompt (zero-shot CoT) achieves surprisingly strong results without any examples. A must-read companion to the Wei et al. paper for understanding the minimal effective dose of reasoning prompts.

 PAPER

[Wang, X. et al. \(2023\). Self-Consistency Improves Chain of Thought Reasoning in Language Models. ICLR 2023.](#)

Introduces self-consistency: sampling multiple reasoning paths and taking the majority vote. Achieves significant accuracy gains over single-path CoT across diverse benchmarks, at the cost of increased API calls.

 PAPER

[Yao, S. et al. \(2023\). Tree of Thoughts: Deliberate Problem Solving with Large Language Models. NeurIPS 2023.](#)

Extends CoT into a tree-structured search where the model explores multiple reasoning branches, evaluates partial solutions, and backtracks. Most effective for problems with discrete solution steps like puzzles and planning tasks.

 PAPER

[Zheng, H. S. et al. \(2023\). Take a Step Back: Evoking Reasoning via Abstraction in Large Language Models.](#)

Proposes step-back prompting, where the model first identifies high-level principles before solving a specific problem. Particularly effective for science and knowledge-intensive reasoning tasks.

 PAPER

[Yao, S. et al. \(2023\). ReAct: Synergizing Reasoning and Acting in Language Models. ICLR 2023.](#)

Introduces the ReAct framework that interleaves reasoning traces with tool-use actions. The foundation for most modern agent architectures, combining CoT's reasoning benefits with the ability to retrieve information and execute code.

 PAPER